

Information Retrieval in High Dimensional Data
Assignment #3, 19.01.2023

Due date: 19.02.2023 at 10pm

Please hand in your solutions via Moodle. Use the attached Jupyter notebook.

Solutions must be handed in by groups. Please state the names of your group members at a prominent place in your submission. (For example, at the beginning of your provided notebook or in a separate text file.)

The Kernel Trick

Task 1: [25 points] On Moodle you will find a Jupyter-Notebook that contains a function for dimensionality reduction via PCA. The function `linear_pca` expects a data matrix $\mathbf{X} \in \mathbb{R}^{p \times N}$ and a number of PCs k and returns the first k PCA scores for the matrix \mathbf{X} .

- Provide code that tests the function with selected images from the provided MNIST training dataset by visualizing the first 2 scores in a scatter plot.
- Complete the function `gram_pca` such that it has the same functionality as `linear_pca` but expects a gram matrix $\mathbf{K} = \mathbf{X}^\top \mathbf{X}$ instead of the data matrix \mathbf{X} as its input. Do not assume that \mathbf{K} was produced from centered data. Note: It is important to be consistent in notation here. E.g., for a data matrix of 1000 MNIST images, we have $\mathbf{X} \in \mathbb{R}^{784 \times 1000}$ and $\mathbf{K} \in \mathbb{R}^{1000 \times 1000}$.
- Test your implementation and show that `gram_pca(dot(X.T,X), k)` yields results equivalent to those of `linear_pca(X, k)`.
- There is as an unknown vector space \mathbb{H} , equipped with an inner product $\langle \cdot, \cdot \rangle_{\mathbb{H}}$ and a function

$$\varphi : \mathbb{R}^p \rightarrow \mathbb{H},$$

such that

$$\langle \varphi(\mathbf{x}), \varphi(\mathbf{y}) \rangle_{\mathbb{H}} = \exp \left(-\frac{\|\mathbf{x} - \mathbf{y}\|^2}{2\sigma^2} \right)$$

holds for every $\mathbf{x}, \mathbf{y} \in \mathbb{R}^p$. The expression on the right-hand side of the equation is called the *Gaussian kernel* and σ is a parameter to choose by hand.

The function `gaussian_kernel_pca` expects a data matrix \mathbf{X} , a reduced dimension number k and a parameter σ . It returns the first k *Kernel PCA* scores of the data. In other words, the function returns the first k PCA scores of

$$\varphi(\mathbf{x}_1), \varphi(\mathbf{x}_2), \dots, \varphi(\mathbf{x}_N),$$

where \mathbf{x}_i denotes the i -th data sample/ i -th column of the data matrix. The function `gaussian_kernel_pca` is already written, but for it to work, the function `compute_gaussian_gram_matrix` must return correct results. Complete `compute_gaussian_gram_matrix` accordingly.

- Test `gaussian_kernel_pca` with some MNIST train images and $\sigma = 1000$.